

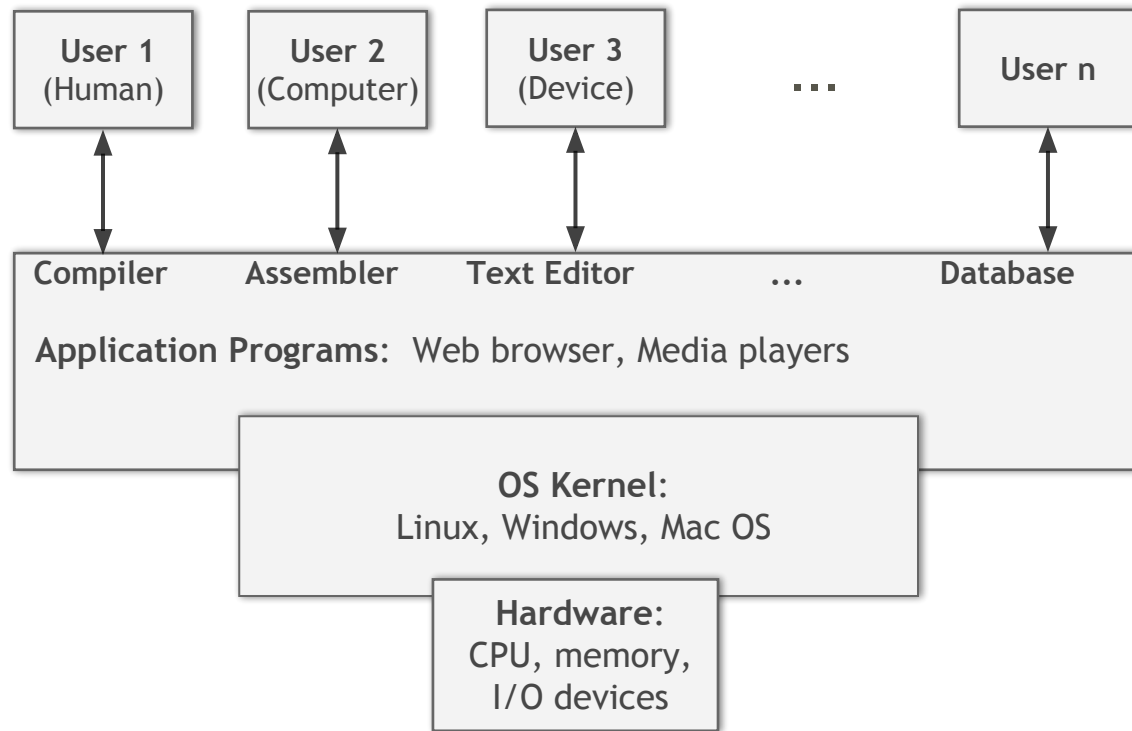
CPSC 457

System calls

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

- provides services to applications, eg.
 - access to hardware, often via higher level abstractions
 - resource management
- these OS services are accessible through **system calls**, aka kernel calls
 - usually via **traps** - software interrupts

Kernel vs. user mode



User mode - no access to hardware, limited instruction set

Kernel mode - full access to hardware, full instruction set

System calls

- when an application wants to access a service / resource of the system:
 - the application must make an appropriate **system call** - call a routine provided by the OS
 - usually done through a mechanism called **trap**
 - trap is a special CPU instruction
 - switches from **user mode** to **kernel mode**
 - invokes a pre-defined trap handler, registered by kernel
 - inside trap handler:
 - OS saves application state
 - OS does the requested operation, eg. involving some hardware
 - OS switches back to user mode and restores application state
 - the application resumes
- from application's perspective, making a system call is just like calling any other routine

- system calls provide an interface to the services made available by OS
 - think of it as an API provided by the OS
 - the interface for system calls varies from OS to OS, although the underlying concepts tend to be similar
 - OS often needs to execute 1000s of system calls per second

Example: copy file

Source File

Destination File

```
Acquire input file name
  Write prompt to screen
  Accept input
Acquire output file name
  Write prompt to screen
  Accept input
Open the input file
If input file doesn't exist, abort
Create output file
If file could not be created, abort
```

```
Loop
  Read byte(s) from input file
  Write byte(s) to output file
Until read or write fails
Close input file
Close output file
Write completion message to screen
Terminate normally
```

even the most simple programs make many system calls

Libraries and system calls

- system calls are minimalistic, and not very easy to use
 - usually implemented in assembly, optimized for performance
 - system call number and parameters usually passed in registers

```
mov  eax,4      ; system call # (sys_write)
mov  ebx,1      ; fd = stdout
mov  edx,4      ; message length
mov  ecx,msg    ; ptr to message
int  0x80      ; trap
```

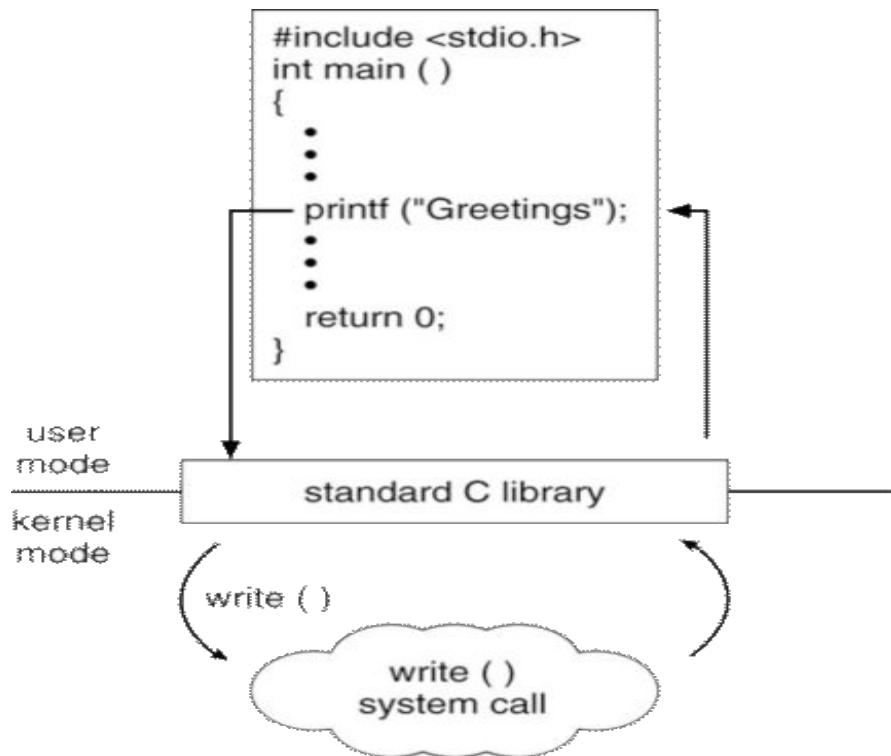
- http://blog.rchapman.org/posts/Linux_System_Call_Table_for_x86_64/
 - quite inconvenient to use from higher level languages
- preferred way to make system calls is through higher-level wrappers
- on Unix-like systems:
 - libc** - a C library, **libstdc++** or **libc++** for C++

Libraries and system calls

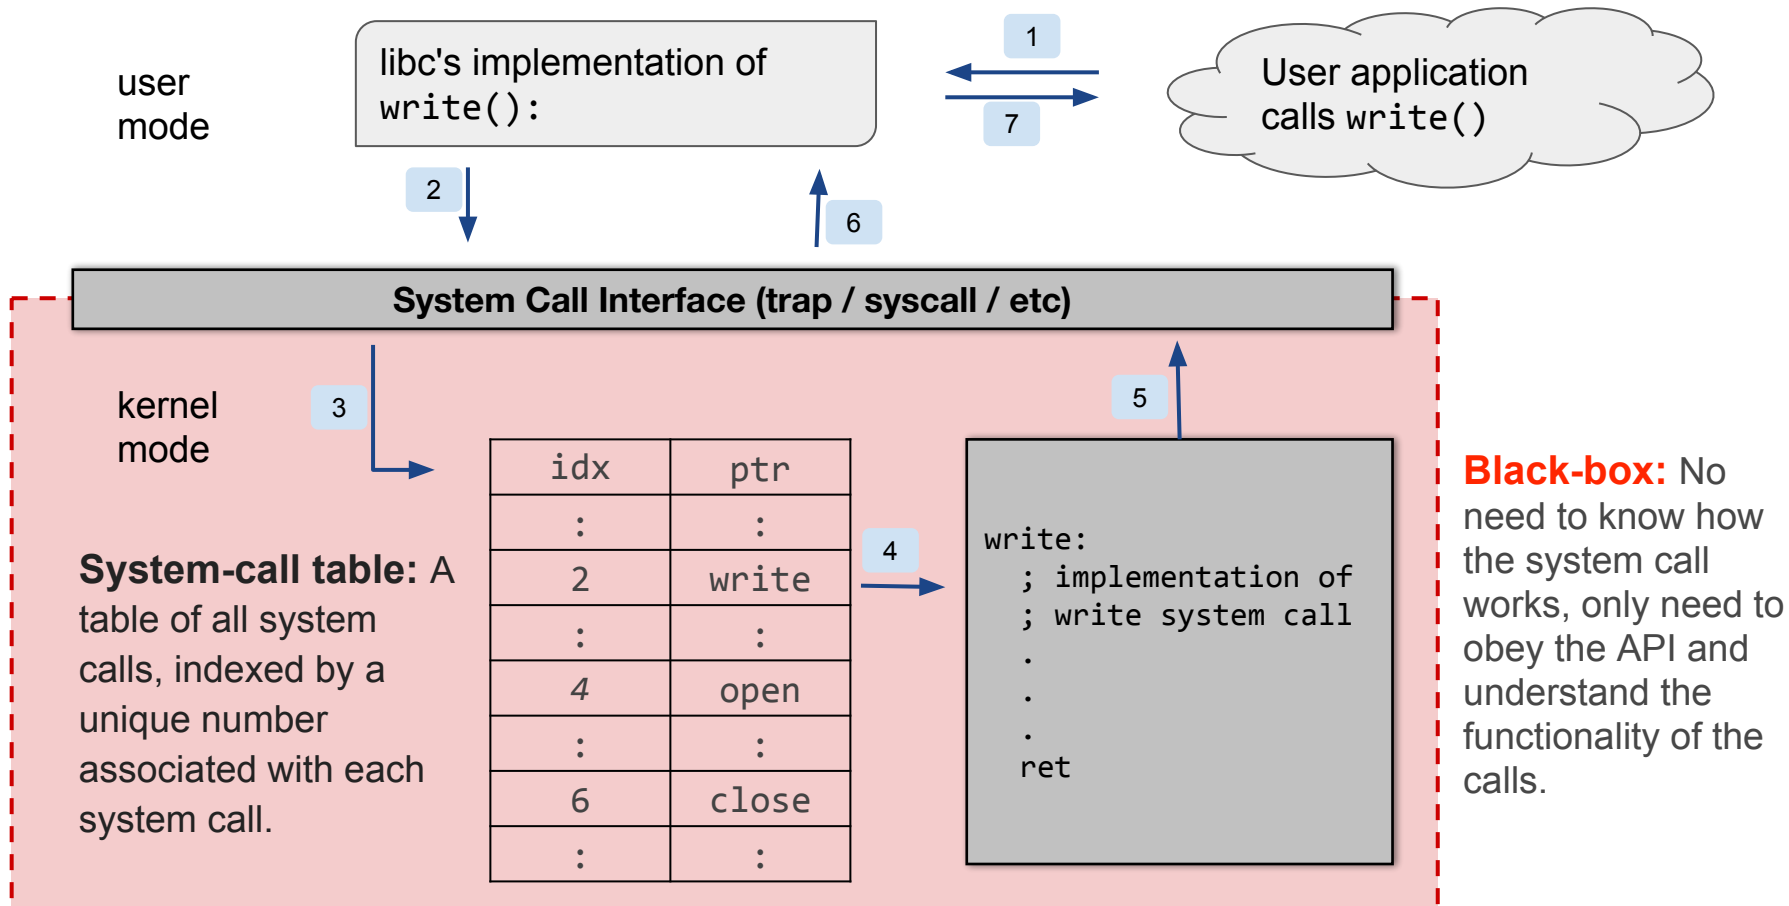
- a library can provide a set of functions (APIs) that are available to an application programmer, including the parameters and the return values
- these APIs often hide the implementation details of system calls
- making system calls via wrappers is more convenient
- an application can compile and run on any system that supports the same API
- added benefit - if the system call ever changes / is deprecated, the program using the wrapper could still continue to function properly
- some common APIs:
 - POSIX APIs for Unix, Linux, Mac OS X
 - Win32 APIs for windows
 - Java APIs for Java virtual machine
- often there is a strong correlation between a function in the API and its associated system call within the kernel, but API != system call

Example: printf(...)

- standard C library provides access to many OS system calls
- for example: `write()`
 - `write()` prepares arguments in registers
 - `write()` calls the `write` system call
 - `write()` takes the value returned by `write` and passes it back to the caller
- but also many useful higher-level APIs, eg. `printf()`
 - `printf()` does some formatting and then calls the system call `write`



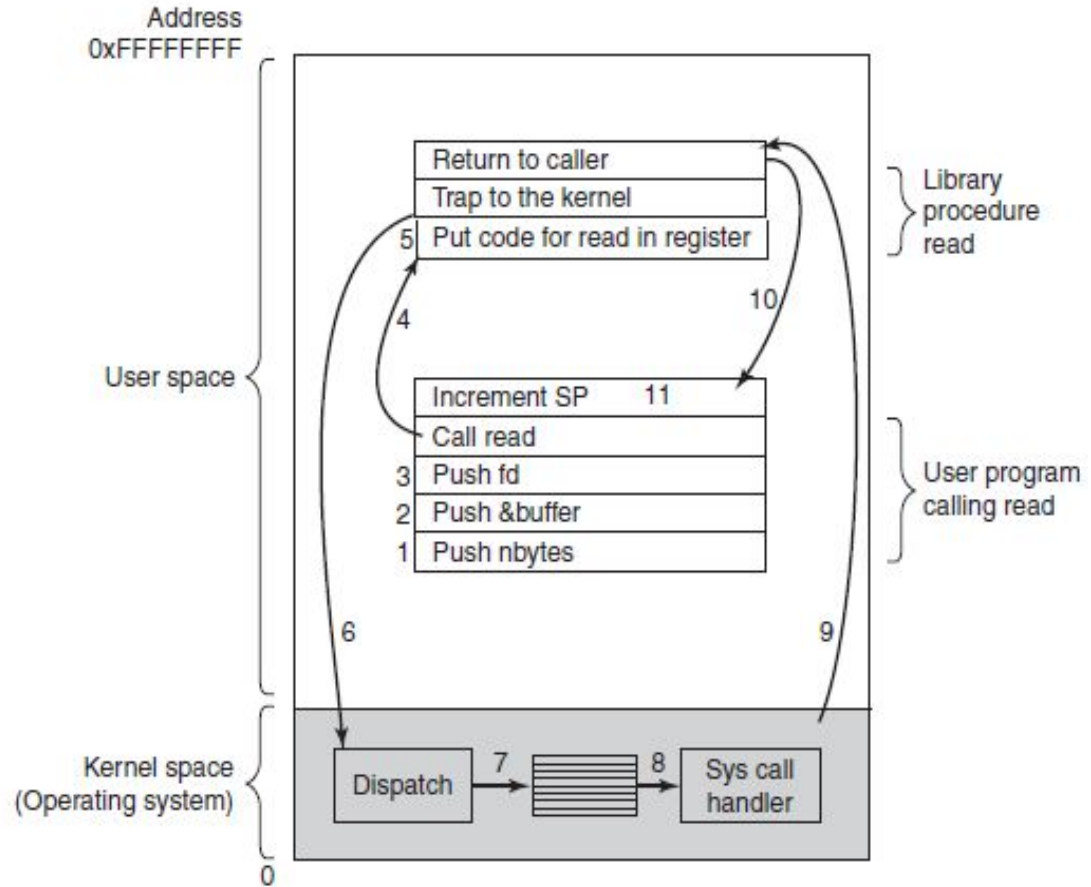
API / System calls / OS relationship



Example: read()

Steps in making a wrapper call

`read(fd, buffer, nbytes)`



Examples of system call APIs in C

File management

| Call | Description |
|---|---|
| <code>fd = open(file, how, ...)</code> | Open a file for reading, writing, or both |
| <code>s = close(fd)</code> | Close an open file |
| <code>n = read(fd, buffer, nbytes)</code> | Read data from a file into a buffer |
| <code>n = write(fd, buffer, nbytes)</code> | Write data from a buffer into a file |
| <code>position = lseek(fd, offset, whence)</code> | Move the file pointer |
| <code>s = stat(name, &buf)</code> | Get a file's status information |

More examples of system call APIs in C

Directory and file system management

| Call | Description |
|---|--|
| <code>s = mkdir(name, mode)</code> | Create a new directory |
| <code>s = rmdir(name)</code> | Remove an empty directory |
| <code>s = link(name1, name2)</code> | Create a new entry, name2, pointing to name1 |
| <code>s = unlink(name)</code> | Remove a directory entry |
| <code>s = mount(special, name, flag)</code> | Mount a file system |
| <code>s = umount(special)</code> | Unmount a file system |

Even more examples of system call APIs in C

Miscellaneous

| Call | Description |
|---|---|
| <code>s = chdir(dirname)</code> | Change the working directory |
| <code>s = chmod(name, mode)</code> | Change a file's protection bits |
| <code>s = kill(pid, signal)</code> | Send a signal to a process |
| <code>seconds = time(&seconds)</code> | Get the elapsed time since Jan. 1, 1970 |

System calls examples (UNIX vs Win32)

| UNIX | Win32 | Description |
|---------|---------------------|--|
| fork | CreateProcess | Create a new process |
| waitpid | WaitForSingleObject | Can wait for a process to exit |
| execve | (none) | CreateProcess = fork + execve |
| exit | ExitProcess | Terminate execution |
| open | CreateFile | Create a file or open an existing file |
| close | CloseHandle | Close a file |
| read | ReadFile | Read data from a file |
| write | WriteFile | Write data to a file |
| lseek | SetFilePointer | Move the file pointer |
| stat | GetFileAttributesEx | Get various file attributes |
| mkdir | CreateDirectory | Create a new directory |

Tracing system calls

- tracing system calls = running an application and logging all system calls
- usually for debugging or performance optimization purposes
- on Linux: `strace`
- on Solaris: `truss`
- on Mac OS X: `dtruss`
- on Windows: Windows Performance Analysis Tools
<https://msdn.microsoft.com/en-us/windows/hardware/commercialize/test/wpt/windows-performance-analyzer>
- note: the same program/command could invoke different set of system calls on different OSes
- refer to the man page for further detail on these commands

Man pages

```
$ man strace
```

Man pages

STRACE(1)

General Commands Manual

STRACE(1)

NAME

strace - trace system calls and signals

SYNOPSIS

```
strace [-CdffhikqrrttTvVxxy] [-In] [-bexecve] [-eexpr]... [-acolumn]
[-ofile] [-ssstrsize] [-Ppath]... -ppid... / [-D] [-Evar[=val]]...
[-username] command [args]
```

```
strace -c[df] [-In] [-bexecve] [-eexpr]... [-Ooverhead] [-Ssortby]
-ppid... / [-D] [-Evar[=val]]... [-username] command [args]
```

DESCRIPTION

In the simplest case strace runs the specified command until it exits. It intercepts and records the system calls which are called by a process and the signals which are received by a process. The name of each system call, its arguments and its return value are printed on standard error or to the file specified with the -o option.

Strace

```
$ strace cat sample.txt
```

```
$ strace ./readFile sample.txt
```

```
$ strace -c cat sample.txt
```

```
$ strace -c ./readFile sample.txt
```

Strace

```
$ strace cat sample.txt

...
open("readme.txt", O_RDONLY)           = 3
fstat(3, {st_mode=S_IFREG|0600, st_size=4, ...}) = 0
fadvise64(3, 0, 0, POSIX_FADV_SEQUENTIAL) = 0
mmap(NULL, 1056768, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fd581f6e000
read(3, "hey\n", 1048576)              = 4
write(1, "hey\n", 4hey
)                                       = 4
read(3, "", 1048576)                 = 0
munmap(0x7fd581f6e000, 1056768)      = 0
close(3)                             = 0
close(1)                             = 0
close(2)                             = 0
exit_group(0)                        = ?
...
```

Strace


```
$ strace -c cat sample.txt
```

```
...
```

| % time | seconds | usecs/call | calls | errors | syscall |
|--------|----------|------------|-------|--------|------------|
| 35.27 | 0.000073 | 18 | 4 | | open |
| 16.43 | 0.000034 | 3 | 10 | | mmap |
| 8.21 | 0.000017 | 4 | 4 | | mprotect |
| 8.21 | 0.000017 | 9 | 2 | | munmap |
| 7.73 | 0.000016 | 3 | 5 | | fstat |
| 4.83 | 0.000010 | 2 | 6 | | close |
| 4.35 | 0.000009 | 3 | 3 | | read |
| 3.86 | 0.000008 | 8 | 1 | | write |
| 3.86 | 0.000008 | 8 | 1 | 1 | access |
| 3.38 | 0.000007 | 2 | 4 | | brk |
| 1.93 | 0.000004 | 4 | 1 | | execve |
| 0.97 | 0.000002 | 2 | 1 | | arch_prctl |
| 0.97 | 0.000002 | 2 | 1 | | fadvise64 |
| 100.00 | 0.000207 | | 43 | 1 | total |

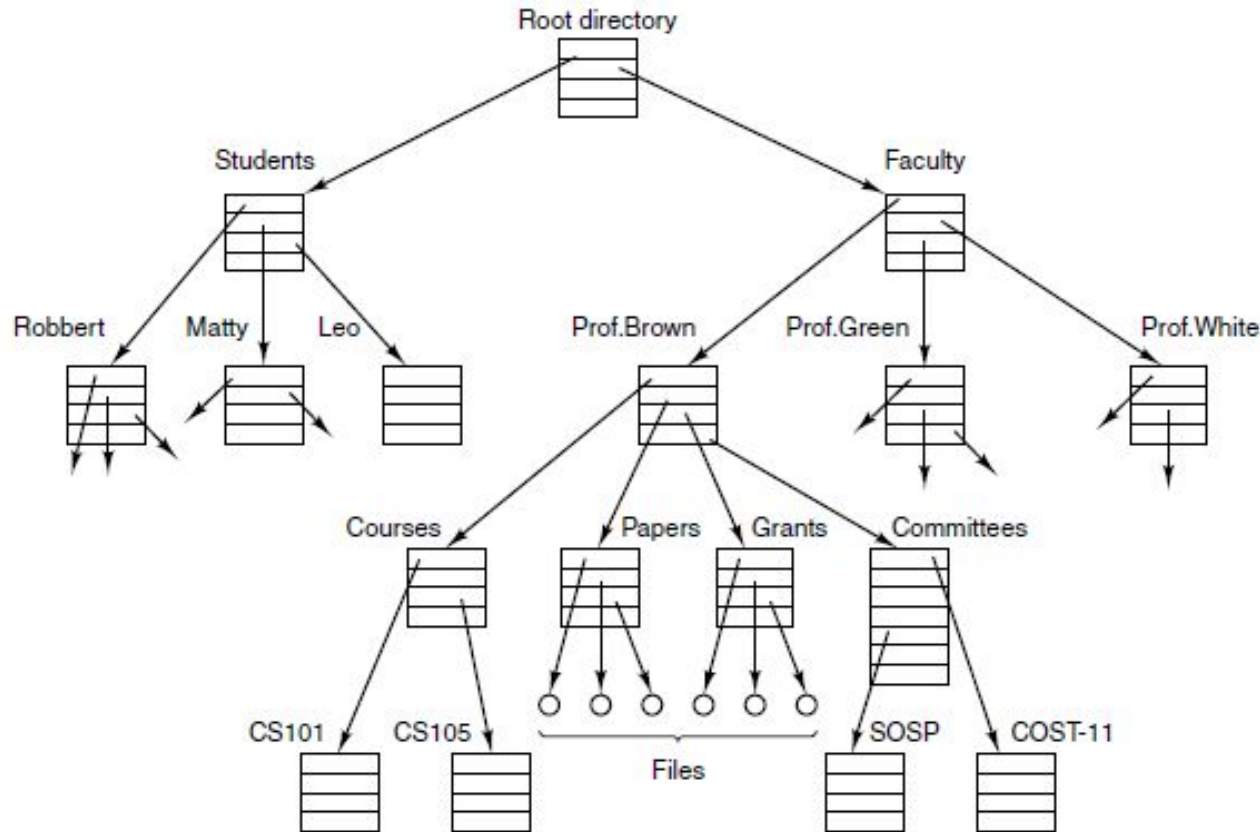
- key concept in all operating systems
- quick definition: a program in execution
- process is associated with
 - an address space
 - set of resources
 - program counter, stack pointer
 - unique identifier (process ID)
 - ... anything else?
- process can be thought of as a container that holds all information needed by an OS to run a program

Process tree




- processes are allowed to create new processes
- A creates two **child processes**: B and C
- B creates three child processes: D, E and F
- A is the **parent process** of B
- B is a parent process of E


File system - tree structure (subdirectories and files)



File system mounting




(a)



(b)

- on most Unix filesystems, other filesystems are mounted into an existing filesystems at an arbitrary location (subdirectory)
- (a) before mounting
- (b) after mounting



- on unix systems, two processes can communicate with each other via a pipe
- pipes are accessed using file I/O APIs

```
$ ls -altr | tail -10
```

- UNIX-like OSs make use of files and associated APIs for different operations / services
- pipes - interprocess communication
- sockets - networking
- devices (`/dev`)
 - block devices - disks
 - character devices - terminals
- random number generator (`/dev/random`)
- export kernel parameters (`/proc` and `/sys`)
 - pseudo filesystems containing virtual files
 - eg. information about processes, memory usage, hardware devices

Questions?