

# CPSC 457

## Processes - part 1

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

- we want the ability to run multiple programs at the same time (multitasking)
- for this we need firm control and compartmentalization of the various programs
- to this end we create an abstraction of **program in execution** and call it a **process**

- we need the process abstraction to implement **multitasking**
- multitasking allows an illusion of parallelism
  - running  $N$  processes with  $M$  CPUs, while  $N > M$
  - works even with a single CPU:
    - run program (i) for a fraction of a second
    - switch to program (i+1)
    - repeat
- multitasking allows us to reduce CPU idling during I/O
  - CPU could be given to another process rather than remain idle
- multitasking is only practical when memory is big enough to hold multiple running programs

# Program != Process

---


- a program is a **passive** entity
  - an executable file containing a list of instructions, usually stored on disk
- a process is an **active** entity
  - associated with a program counter and other resources
- a program **becomes** a process when it is loaded into memory for execution
- a single program can be used to create multiple processes
  - eg. running multiple instances of one executable

# Analogy: baking and OS

Baking	OS
recipe	○ program
the chef	○ CPU
flour, eggs, sugar	○ input data
reading the recipe, getting the ingredients, baking the cake	○ process
phone call	○ interrupt

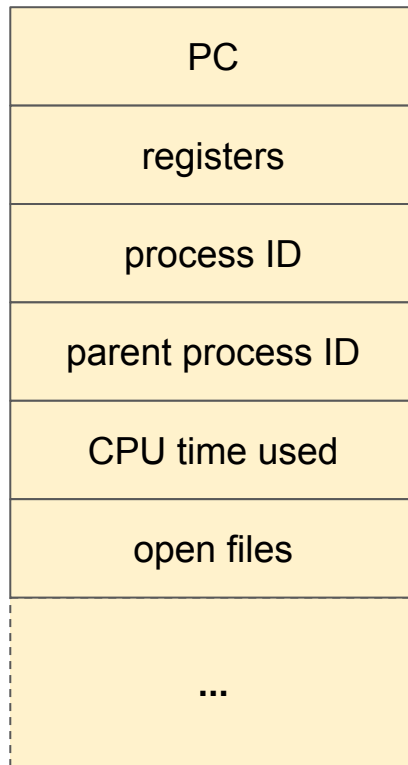
# A process in memory

- each process gets its own **address space**
  - part of memory available to a process, decided by OS
  - on modern OSes it is a *virtual* address space (0 - max), isolated from other processes
- **text section**: the program code
- **data section**: global variables, constant variables
- **heap**: memory for dynamic allocation during runtime
- **stack**: temporary data (parameters, return address, local variables)
- plus many other bits of information needed by the OS for management, usually grouped in a Process Control Block data structure



# Process control block (PCB)

- each process is represented in the OS by a PCB that includes:
  - process state
  - program counter
  - CPU registers
  - CPU-scheduling info priority, pointers to the queue, other parameters
  - memory management info: page tables, segment tables, etc.
  - accounting info: CPU time, timeout values, process numbers, etc.
  - I/O status info: open files, I/O devices, etc.
- the **process table** is a collection of all PCBs



# Some of the fields of a PCB

## Process management

program counter  
registers  
stack pointer  
process state  
priority  
scheduling parameters  
process ID  
parent process  
process group  
signals  
process start time  
CPU time used  
children's CPU time used  
time of next alarm  
...

## Memory management

pointer to text segment  
pointer to data segment  
pointer to stack segment  
...

## File management

root directory  
working directory  
file descriptors  
user ID  
group ID  
...

Look for “**task\_struct**” in Linux kernel sources

<https://github.com/torvalds/linux/blob/master/include/linux/sched.h>



# Operations on processes

- processes need to be created and deleted dynamically and OS must provide mechanisms for this
- process creation, eg. `fork()` in UNIX
  - **parent process** - the process that is creating a new process
  - **child process (child)** - the newly created process
  - processes in the system form a **process tree**
  - each process gets **PID** - a unique **process identifier**
- process execution, eg. `fork()`
- process termination, eg. `exit()` or `kill()`
  - to let the OS delete the process
  - termination can be (typically) only requested by the process or its parent
- other operations: synchronization, communication, ...

# A multiprocess program in C

```
$ man fork
```

```
pid_t fork(void);
```

`fork()` creates a new process by duplicating the calling process. The new process is referred to as the child process. The calling process is referred to as the parent process.

The child process and the parent process run in separate memory spaces. At the time of `fork()` both memory spaces have the same content. Memory writes, file mappings (`mmap(2)`), and unmappings (`munmap(2)`) performed by one of the processes do not affect the other.

The child process is an exact duplicate of the parent process except for the following points:

...

# A multiprocess program in C

```
#include <unistd.h>

int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

Possible output:

???

# A multiprocess program in C

```
#include <unistd.h>
```

```
int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

Possible output:

Hello  
world.  
world.

Are other outputs  
possible?

# A multiprocess program in C

```
#include <unistd.h>
```

```
int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

Possible output:

Hello  
world.  
world.

Another possible output:

Hello  
worworld.  
rld.

Are other outputs  
possible?

# A multiprocess program in C

```
#include <unistd.h>
```

```
int main()
{
    printf("Hello\n");
    /* create & run child process - a duplicate of parent */
    fork();
    /* both parent and child will execute the next line */
    printf("world.\n");
}
```

If `fork()` fails, it returns **-1**. You should always check the return value of a system call.

Possible output:

Hello  
world.  
world.

Another possible output:

Hello  
worworld.  
rld.

Output if `fork()` fails:

Hello  
world.

# A multiprocess program in C

```
int main()
{
    /* create & run child process - a duplicate of parent
     * and remember the return value */
    pid_t pid = fork();
    /* both parent and child will execute the next line,
     * but will have different value for pid:
     * 0 for child, non-zero for parent */
    printf("My pid is %d.\n", pid);
}
```

Possible output:

My pid is 7.  
My pid is 0.

Possible output:

My pid is 0.  
My pid is 7.

Possible output:

My pid is -1.

# Homework - can you predict the output?

```
int main()
{
    fprintf( stderr, "A\n");
    fork();
    fprintf( stderr, "B\n");
    fork();
    fprintf( stderr, "C\n");
}
```

```
int main()
{
    for(int i=0 ; i<4 ; i++ ) {
        fork();
    }
    printf("X");
}
```



# Fork bomb

```
int main()
{
    while(1) {
        fork();
    }
    printf("X");
}
```

**WARNING:** please do not run this on CPSC servers.

# A multiprocess program in C

```
$ man execl
```

```
int execl(const char *path, const char *arg, ...);  
int execlp(const char *file, const char *arg, ...);  
...
```

The `exec()` family of functions replaces the current process image with a new process image. The functions described in this manual page are front-ends for `execve(2)`. (See the manual page for `execve(2)` for further details about the replacement of the current process image.)

The initial argument for these functions is the name of a file that is to be executed.

```
...
```

# A multiprocess program in C

```
$ man execve
```

```
int execve(const char *filename, char *const argv[], char *const envp[]);
```

`execve()` executes the program pointed to by filename. filename must be either a binary executable, or a script starting with a line of the form:

```
#! interpreter [optional-arg]
```

For details of the latter case, see "Interpreter scripts" below.

argv is an array of argument strings passed to the new program. By convention, the first of these strings should contain the filename associated with the file being executed. envp is an array of strings, conventionally of the form key=value, which are passed as environment to the new program. Both argv and envp must be terminated by a null pointer.

...

...

# A multiprocess program in C

```
int main()
{
    pid_t  pid;

    pid = fork(); /* create & run child process - a duplicate of parent */
    /* both parent and child will execute the next line */
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", "-l", NULL); /* replace process with 'ls -l' */
    }
    else { /* parent process will wait for the child to complete */
        printf("Waiting for child process %d\n", pid);
        wait(NULL);
        exit(0);
    }
}
```

# A multiprocess program in C

```
#include <stdlib.h>
```

```
int main()
{
    system("/bin/ls");
}
```

```
$ man system
```

```
...
```

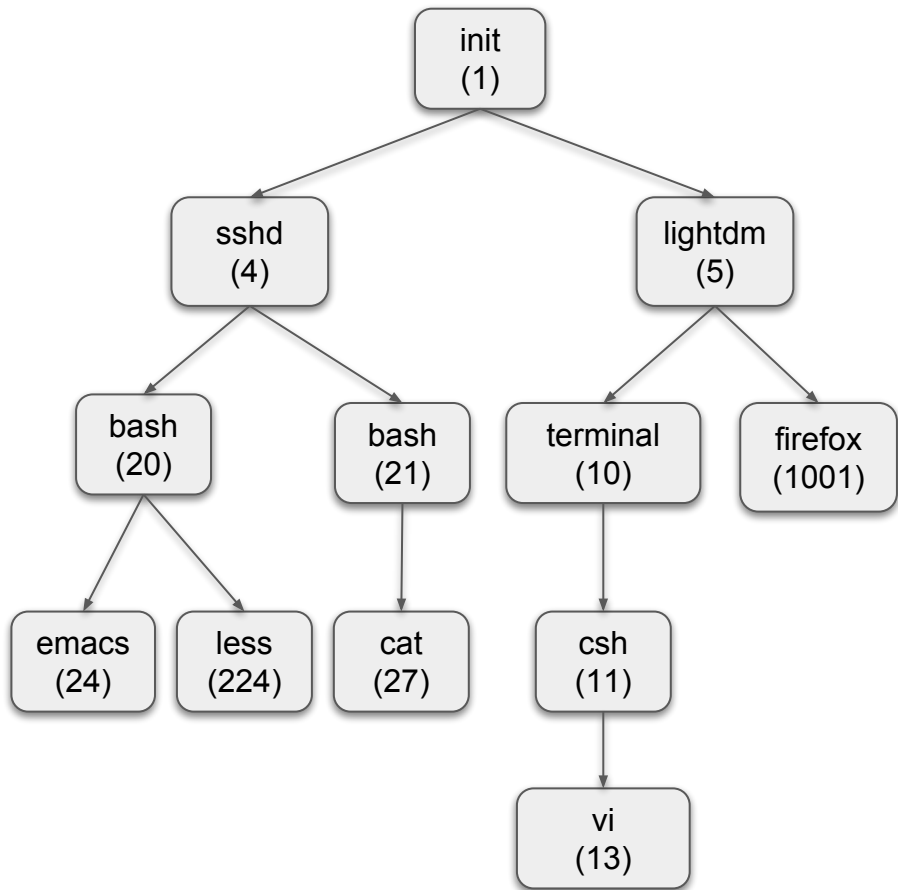
The `system()` library function uses `fork(2)` to create a child process that executes the shell command specified in `command` using `execl(3)` as follows:

```
execl("/bin/sh", "sh", "-c", command,
      (char *) 0);
```

`system()` returns after the command has been completed.

```
...
```

# Process tree



- **parent process** - the creating process
- **child process (child)** - the newly created process
- **PID** - the unique process identifier for each process
- in Unix, parent and child processes continue to be associated, forming a process hierarchy
- in Windows, all processes are equal, the parent process can give the control of its children to any other process

# init process

- `init` or `systemd` is the first process started after booting
  - older UNIX systems used `init` based systems
  - many popular Linux systems now switched to `systemd`
- `init` is the **ancestor** of all user processes (direct or indirect parent), i.e. root of process tree
- `init` always has PID = 1
- **orphaned** processes are adopted by `init`
- printing a process tree

```
$ pstree
$ ps axjf
```
- note: some special system 'processes' are created by kernel during bootstrap, and do not have to be descendants of `init`, such as `swapper` and `pagedaemon`

- Which one of the following executes in kernel mode?
  - ☐ A user program
  - ☐ A library function call
  - ☐ A system call
  - ☐ A system call wrapper function
  
- In C, `printf()` is a system call.
  - ☐ True
  - ☐ False



- When 4 programs are executing on a computer with a single CPU, how many program counters are there?
- When does a program become a process?

- What is the name of the PCB data structure in Linux?
- Name some of fields in a PCB.
- On UNIX systems, what is the name of the process that is the ancestor of all user processes?

# Summary

---

- system calls + related APIs
- processes
- processes implementation
  - PCB, address space
- fork, exec

Reference: 1.5, 1.6, 2.1.1, 2.1.2, 2.1.6 (Modern Operating Systems)

2.1 - 2.4, 3.1, 3.3 (Operating System Concepts)

# Questions?