

CPSC 457

Concurrent programming

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

- threads and fork()
- thread cancellation
- race conditions
- critical sections

Threads and `fork()`

- is it ok to call `fork()` in a program with multiple threads?
 - what should happen?
 - what does happen?
- what actually happens:
 - only the calling thread survives, other threads are not duplicated
 - this creates a problem if synchronization mechanisms were used
 - it's possible to register a callback in case `fork()` is called using `pthread_atfork()`
- general advice: avoid using `fork()` in programs with multiple threads
- some usages are safe, eg.:
 - `fork()` is immediately followed by `execve()` to execute external program, or
 - `fork()` is executed before creating any threads

Thread cancellation

- scenario: parallelizing database search
- multiple threads searching different parts of the database
- one thread finds the result
- how do we notify the other threads to stop searching?

- two general approaches:
 - asynchronous cancellation
 - deferred cancellation (aka. synchronous cancellation)

Asynchronous thread cancellation

- one thread manually terminates the target thread
 - `pthread_kill(threadid, SIGUSR1)`
 - target thread is killed instantly*
- common problem: what happens to the data currently updated by the thread that is killed?
 - killed thread has no chance to "clean up"
 - this can (likely) lead to leaving data in undefined state
- usually a much better solution is to use synchronous thread cancellation

Deferred / Synchronous thread cancellation

- controlling thread somehow *indicates* it wishes to cancel a thread
 - many options
 - eg. by setting some shared global flag
 - or use `pthread_cancel()` and related mechanisms (`man pthread_cancel` for details)
- target thread periodically checks whether it should terminate
 - checking done only at **cancellation points** at which it can be canceled safely
- some issues:
 - less performance
 - checking the cancellation flag requires at least 1 instruction...
 - target thread will not react immediately
 - may run for a while after cancellation requested
 - eg. could continue to report results for a while after cancellation request issued
- more flexible than asynchronous cancellation, but also more complex to implement

Deferred cancellation example

```
void * thread_print(void * tid) {
    while ( 1) {
        printf("thread %ld running\n", tid);
        sleep( 1);
    }
    pthread_exit(0); // not necessary
}

int main() {
    pthread_t threads[NUMBER_OF_THREADS];
    for (long i = 0; i < NUMBER_OF_THREADS; i++) {
        if( 0 != pthread_create(& threads[i], NULL, thread_print, (void *) i)) {
            printf("Oops, pthread_create failed\n");
            exit(-1);
        }
    }
    sleep(5); // pretend to do something
    ???
    for (long i = 0; i < NUMBER_OF_THREADS; i++)
        pthread_join(threads[i], NULL);
    exit(0);
}
```

keep printing
message forever

here we need to
cancel the threads,
otherwise this
program will also
run forever

Deferred cancellation example

```

volatile int cancel_flag = 0;
void * thread_print(void * tid) {
    while(1) {
        printf("thread %ld running\n", tid);
        sleep(1);
        if(cancel_flag ) return;
    }
}

int main() {
    pthread_t threads[NUMBER_OF_THREADS];
    for (long i = 0; i < NUMBER_OF_THREADS; i++) {
        if( 0 != pthread_create(& threads[i], NULL, thread_print, (void *) i)) {
            printf("Oops, pthread_create failed.\n"); exit(-1);
        }
    }
    sleep(5); // pretend to do something
    cancel_flag = 1;
    for (long i = 0; i < NUMBER_OF_THREADS; i++)
        pthread_join(threads[i], NULL);
    exit(0);
}

```

global flag

periodically check
the global flag
(cancellation point)

set global flag to
request cancellation

works fine on x86, but
should use an
atomic type / operations
eg. `std::atomic<bool>`
for portability

Race conditions

```
// global variable counter  
int counter;
```

```
void incr() {  
    int x = counter;  
    x = x + 1;  
    counter = x;  
}
```

```
int main() {  
    counter = 0;  
    incr();  
    incr();  
    printf( "%d\n", counter);  
}
```

Output:

2

... every time

Race conditions

```
// global variable counter  
int counter;
```

```
void incr() {  
    int x = counter;  
    x = x + 1;  
    counter = x;  
}
```

```
int main() {  
    counter = 0;  
    pthread_create(..., incr);  
    pthread_create(..., incr);  
    pthread_join ...  
    printf( "%d\n", counter);  
}
```

Thread 1:

```
void incr() {  
    int x = counter;  
    x = x + 1;  
    counter = x;  
}
```

Thread 2:

```
void incr() {  
    int x = counter;  
    x = x + 1;  
    counter = x;  
}
```

What is the value in **counter** after
both threads finish executing
incr()?

Race conditions

```
// global variable counter
int counter;

void incr() {
    int x = counter;
    x = x + 1;
    counter = x;
}

int main() {
    counter = 0;
    pthread_create(..., incr);
    pthread_create(..., incr);
    pthread_join ...
    printf( "%d\n", counter);
}
```

Thread 1	Thread 2	counter
		0
x = counter;		0
x = x + 1;		0
counter = x;		1
	x = counter;	1
	x = x + 1;	1
	counter = x;	2

possible **execution sequence**
leading to counter = 2

Race conditions

```
// global variable counter
int counter;

void incr() {
    int x = counter;
    x = x + 1;
    counter = x;
}

int main() {
    counter = 0;
    pthread_create(..., incr);
    pthread_create(..., incr);
    pthread_join ...
    printf( "%d\n", counter);
}
```

Thread 1	Thread 2	counter
		0
x = counter;		0
	x = counter;	0
	x = x + 1;	0
	counter = x;	1
x = x + 1;		1
counter = x;		1

another possible execution sequence
leading to **counter = 1 !!!**

Race conditions

```
// global variable counter  
int counter;
```

Would this get rid
of the race
condition?

```
void incr() {  
    int x = counter;  
    x = x + 1;  
    counter = x;  
}
```

```
void incr() {  
  
    counter ++;  
}
```

```
int main() {  
    counter = 0;  
    pthread_create(..., incr);  
    pthread_create(..., incr);  
    pthread_join ...  
    printf( "%d\n", counter);  
}
```

Race conditions

```
int counter;
```

```
int incr1() {  
    int x = counter;  
    x = x + 1;  
    counter = x;  
}
```



```
mov eax, DWORD PTR counter[rip]  
mov DWORD PTR [rbp-4], eax  
add DWORD PTR [rbp-4], 1  
mov eax, DWORD PTR [rbp-4]  
mov DWORD PTR counter[rip], eax
```

```
int incr2() {  
    counter ++;  
}
```



```
mov eax, DWORD PTR counter[rip]  
add eax, 1  
mov DWORD PTR counter[rip], eax
```

To see how GCC compiles your code into assembly instructions, you can try:

```
$ gcc -S -fverbose-asm test.c
```

Or use an online tool, eg: <https://godbolt.org/z/WTPzC2> (full)

Race conditions in software

- **race condition** is a behavior where the output is dependent on the sequence or timing of other uncontrollable events (eg. context switching, scheduling on multiple CPUs)
- often a result of multiple processes/threads operating on a shared state/resource, eg.:
 - modifying shared memory
 - reading/writing to files
 - modifying filesystems
 - reading/writing to databases
- debugging race conditions is not fun
 - many test runs may produce the same output, often correct
<https://repl.it/@pavolfederl/thread-counter-race-condition>
 - then, in a rare situation the output might be different, eg. when system was less/more busy
- we want to avoid race conditions
 - how?

Avoiding race conditions

- we need to prevent more than one process/thread from accessing the shared resource at any given time
- approach:
 - identify **critical sections** in code where this could happen
 - enforce **mutual exclusion** to make sure it does not happen

Critical sections and mutual exclusion

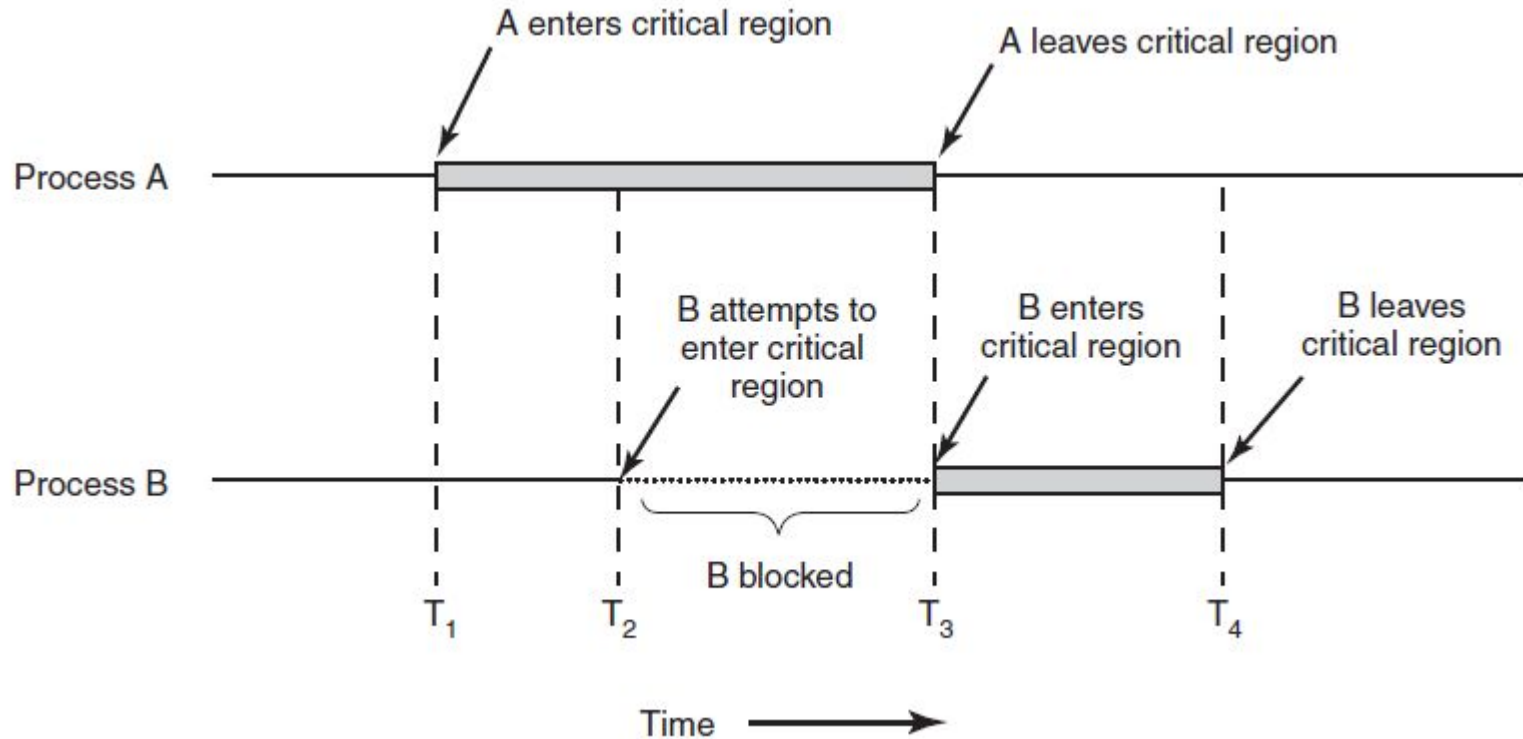
- **critical section** / **critical region**: part of the program that accesses the shared resource in a way that could lead to races or other undefined/unpredictable/unwanted behaviour

```
int counter;                ← shared resource
void incr() {
    int x = counter;
    x = x + 1;
    counter = x;
}
```

← critical section

- if we can arrange tasks such that no two processes or threads will ever be in their critical sections at the same time, we could avoid the race condition (achieving **mutual exclusion**)

Critical sections and mutual exclusion



Critical sections

- requirements to avoid race conditions:
 - no two processes may be simultaneously inside their critical sections
 - no assumptions may be made about the speeds or the number of CPUs
 - no process running outside its critical region may block other processes
 - no process should have to wait forever to enter its critical section

Questions?