

# AN UPDATE IS AVAILABLE FOR YOUR COMPUTER

stickfigures.com

COOL, MORE  
FREE STUFF!



linux

NOT AGAIN!



windows

OOH, ONLY  
\$99!



mac


# CPSC 457

Producer/consumer problem  
Condition variables and Semaphores

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

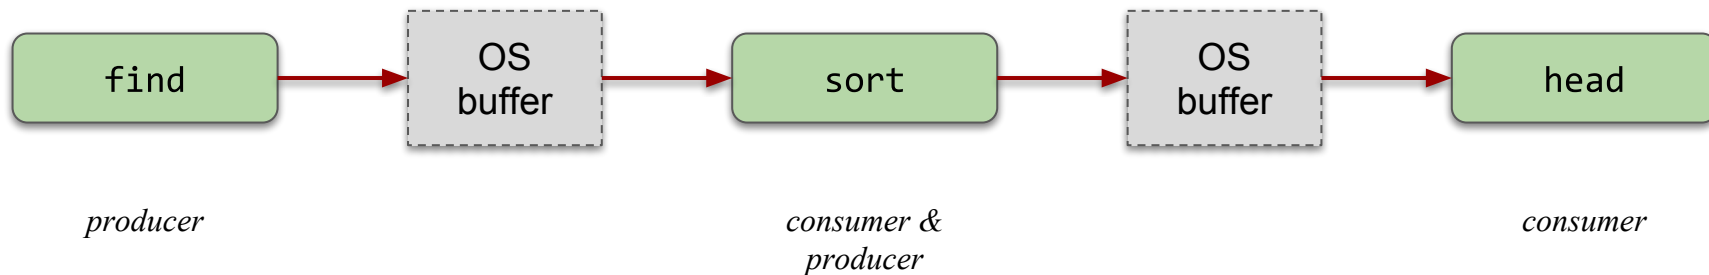
- producer-consumer problem
- mutexes and condition variables
- semaphores

# Producer-consumer problem




# Pipes...

```
$ find . -type f -printf "%-20s%p\n" | sort -nr | head -n 10
```




# The producer-consumer problem

- simplest case: one consumer and one producer processes/threads
- the two processes or threads share a **fixed-size buffer**, used as a queue
- producer puts data into buffer, must wait if buffer full
- consumer takes data out of the buffer, must wait if buffer empty




# Circular buffer

Concept:



Implementation:



- common way to implement fixed-size queue
 

```
typedef struct { ... } item;
item buffer[BUFF_SIZE];
int in = 0; // next free position
int out = 0; // next filled position
int count = 0; // number of items in buffer
```
- buffer is empty when:
 

```
in == out or when count == 0
```
- buffer is full when:
 

```
(in + 1) % BUFF_SIZE == out
```

 or when `count == BUFF_SIZE`

# Possible implementation - naive

## Thread1 – producer thread

```
while(1) {  
    item = produceItem();  
    // wait while buffer is full  
    while((in+1) % BUFF_SIZE == out){}  
    // insert item into buffer  
    buffer[in] = item;  
    in = (in + 1) % BUFF_SIZE;  
    count ++;  
}
```

## Thread 2 – consumer thread

```
while(1) {  
    // wait while buffer is empty  
    while( in == out){}  
    // remove item from buffer  
    item = buffer[out];  
    out = (out+1) % BUFF_SIZE;  
    count --;  
    consumeItem(item);  
}
```

- there is a race condition if we run the above in two different threads
- where?



# Possible implementation - naive

```

Thread1 - producer
    reg1 = count
    reg1 = reg1 + 1
    count = reg1
while(1) {
    item = produceItem();
    // wait while buffer is full
    while((in+1) % BUFF_SIZE == out){}
    // insert item into buffer
    buffer[in] = item;
    in = (in + 1) % BUFF_SIZE;
    count ++;
}
  
```

```

Thread2 - consumer
    reg2 = count
    reg2 = reg2 - 1
    count = reg2
while(1) {
    // wait while buffer is empty
    while( in == out){}
    // remove item from buffer
    item = buffer[out];
    out = (out+1) % BUFF_SIZE;
    count --;
    consumeItem(item);
}
  
```

Possible execution sequence (starting eg. with count=5):

```

T1: reg1 = count           // count=5
T2: reg2 = count           // count=5
T2: reg2 = reg2 - 1        // count=5
T2: count = reg2           // count=4
  
```

```

T1: reg1 = reg1 + 1        // count=4
T1: count = reg1           // count=6
  
```

# Possible implementation with a mutex

## Thread1 – producer thread

```
while(1) {
    item = produceItem();
    // wait while buffer is full
    while((in+1) % BUFF_SIZE == out){}
    // insert item into buffer
    buffer[in] = item;
    in = (in + 1) % BUFF_SIZE;
    pthread_mutex_lock(& mut);
    count ++; // critical section
    pthread_mutex_unlock(& mut);
}
```

## Thread 2 – consumer thread

```
while(1) {
    // wait while buffer is empty
    while( in == out){}
    // remove item from buffer
    item = buffer[out];
    out = (out+1) % BUFF_SIZE;
    pthread_mutex_lock(& mut);
    count --; // critical section
    pthread_mutex_unlock(& mut);
    consumeItem(item);
}
```

- Let's protect the counter using a mutex to remove the race condition with count.
- Can you spot any other problems?

# Possible implementation with a mutex

## Thread1 – producer thread

```
while(1) {  
    item = produceItem();  
    // wait while buffer is full  
    while((in+1) % BUFF_SIZE == out){}  
    // insert item into buffer  
    buffer[in] = item;  
    in = (in + 1) % BUFF_SIZE;  
    pthread_mutex_lock(& mut);  
    count ++;  
    pthread_mutex_unlock(& mut);  
}
```

## Thread 2 – consumer thread

```
while(1) {  
    // wait while buffer is empty  
    while( in == out){}  
    // remove item from buffer  
    item = buffer[out];  
    out = (out+1) % BUFF_SIZE;  
    pthread_mutex_lock(& mut);  
    count --;  
    pthread_mutex_unlock(& mut);  
    consumeItem(item);  
}
```

- this solution would only work for one producer thread and one consumer thread
- what if we wanted multiple producers and/or multiple consumers?
- another important problem: busy wait!

# Possible implementation with a mutex (v2)

## Thread1 – producer thread

```
while(1) {  
    item = produceItem();  
    pthread_mutex_lock(& mut);  
    // wait while buffer is full  
    while((in+1) % BUFF_SIZE == out){}  
    // insert item into buffer  
    buffer[in] = item;  
    in = (in + 1) % BUFF_SIZE;  
    count ++;  
    pthread_mutex_unlock(& mut);  
}
```

## Thread 2 – consumer thread

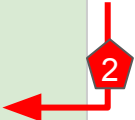
```
while(1) {  
    pthread_mutex_lock(& mut);  
    // wait while buffer is empty  
    while( in == out){}  
    // remove item from buffer  
    item = buffer[out];  
    out = (out+1) % BUFF_SIZE;  
    count --;  
    pthread_mutex_unlock(& mut);  
    consumeItem(item);  
}
```

- let's make one big critical section
- this naive attempt to add support for multiple consumers/producers unfortunately leads to a **deadlock**
- can you find it?

# Possible implementation with a mutex (v2)

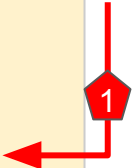
## Thread1 – producer thread

```
while(1) {  
    item = produceItem();  
    pthread_mutex_lock(& mut);  
    // wait while buffer is full  
    while((in+1) % BUFF_SIZE == out){}  
    // insert item into buffer  
    buffer[in] = item;  
    in = (in + 1) % BUFF_SIZE;  
    count ++;  
    pthread_mutex_unlock(& mut);  
}
```



## Thread 2 – consumer thread

```
while(1) {  
    pthread_mutex_lock(& mut);  
    // wait while buffer is empty  
    while( in == out){}  
    // remove item from buffer  
    item = buffer[out];  
    out = (out+1) % BUFF_SIZE;  
    count --;  
    pthread_mutex_unlock(& mut);  
    consumeItem(item);  
}
```



- one possible deadlock: buffer is empty, and consumer enters its critical section ...
  - consumer spins inside CS, producer blocks on trying to acquire mutex
- another deadlock: can you find it?

# Condition variables (CVs)

- condition variables are another type of synchronization primitives
- used together with mutexes
- perfect for implementing critical sections containing loops waiting for some 'condition'

```
// critical section protected with mutex m
lock(m)
...
while( ! some_condition ) {;}
...
unlock(m)
```

- where the condition can only become true if another thread runs its critical section
- CVs make it simple to implement critical sections like the one above

# Condition variables

a common pattern of using CVs:

- a thread enters its critical section (locks a mutex)
- inside CS, thread needs to wait for some condition to become true
- but the condition can only become true by allowing some other thread to lock the mutex
- to facilitate this, the thread puts itself to sleep and simultaneously releases the mutex

```
// critical section protected with mutex m
lock(m)
...
while( ! some_condition ) { wait(cv); }
...
unlock(m)
```

- now some other thread can lock the mutex and execute code that will satisfy the condition

# Condition variables

- eventually some other thread
  - locks the mutex (optional)
  - changes some state that will satisfy the condition
  - notifies the waiting thread via the condition variable
  - releases mutex (optional)

```
// some other thread...  
lock(m)  
...  
some_condition = TRUE;  
signal(cv);  
...  
unlock(m)
```

- the waiting thread then wakes up, and acquires the mutex back automatically



# Condition variables

```
pthread_mutex_t mutex; // mutex
pthread_cond_t cond;   // condition variable
```

- **pthread\_cond\_wait(&cond, &mutex);**
  - atomically releases mutex and causes the calling thread to block, until some other thread calls `pthread_cond_signal(&cond)`
  - after returning, the mutex is automatically re-acquired
  - after returning, the condition must be rechecked !!! (spurious wakeups)
- **pthread\_cond\_signal(&cond);**
  - wakes up one thread waiting on cond
  - if no threads waiting on cond, the signal is lost
  - must be followed by `pthread_mutex_unlock()` if the blocked thread uses the same mutex

# Condition variables

---

- `pthread_cond_init(& cond, & attr)`
  - creates condition variable
- `pthread_cond_destroy(& cond)`
  - destroys a condition variable
- `pthread_cond_broadcast(& cond)`
  - wakes up **all** threads waiting on the condition

# Condition variable example

thread 1:

```
while(1) {  
    pthread_mutex_lock(&mutex);  
    while(count == 0) {  
        pthread_cond_wait(&cond, &mutex);  
    }  
    counter --;  
    pthread_mutex_unlock(&mutex);  
}
```

thread 2:

```
while(1) {  
    pthread_mutex_lock(&mutex);  
    counter ++;  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
}
```

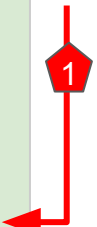
- thread 1 - decrementing counter, but never below 0
- thread 2 - incrementing counter
- no deadlocks
- no busy waiting

# Let's fix this

- deadlock due to one thread stuck in an infinite loop after locking mutex
- while other thread has no chance to run its CS to allow the other thread to exit the loop


## Thread1 – producer thread

```
while(1) {  
    item = produceItem();  
    pthread_mutex_lock(& mut);  
    // wait while buffer is full  
    while((in+1) % BUFF_SIZE == out){}  
    // insert item into buffer  
    buffer[in] = item;  
    in = (in + 1) % BUFF_SIZE;  
    count ++;  
    pthread_mutex_unlock(& mut);  
}
```



## Thread 2 – consumer thread

```
while(1) {  
    pthread_mutex_lock(& mut);  
    // wait while buffer is empty  
    while( in == out){}  
    // remove item from buffer  
    item = buffer[out];  
    out = (out+1) % BUFF_SIZE;  
    count --;  
    pthread_mutex_unlock(& mut);  
    consumeItem(item);  
}
```



# Consumer/producer with condition variables

```
pthread_mutex_t mut;  
pthread_cond_t full, empty;
```

## producer thread:

```
while(1) {  
    item = produceItem();  
    pthread_mutex_lock(& mut);  
    while((in+1) % BUFF_SIZE == out){  
        pthread_cond_wait(& full,& mut);  
    }  
    buffer[in] = item;  
    in = (in + 1) % BUFF_SIZE;  
    count ++;  
    pthread_mutex_signal(& empty);  
    pthread_mutex_unlock(& mut);  
}
```

## consumer thread:

```
while(1) {  
    pthread_mutex_lock(& mut);  
    while( in == out){  
        pthread_cond_wait(& empty,& mut);  
    }  
    item = buffer[out];  
    out = (out+1) % BUFF_SIZE;  
    count --;  
    pthread_mutex_signal(& full);  
    pthread_mutex_unlock(& mut);  
    consumeItem(item);  
}
```

# Semaphore

- another synchronization primitive
- think of semaphore as a special integer variable used for signalling among processes
- the value could indicate number of available units of some resource
- supports three operations:
  - **initialization**
    - can be initialized with any value (0 ... max)
  - **decrement**
    - reduce semaphore by 1,
    - blocks the calling process if going below 0

`down(s)`, `wait(S)` or `sem_wait(S)`
  - **increment**
    - increase value by 1;
    - and possibly unblock another blocked process

`up(S)`, `signal(S)` or `sem_post(S)`

*can be used to  
protect critical  
sections similar to a  
mutex*

```
semaphore s;  
  
wait(s);  
    // Critical Section  
signal(s);
```

- can be used to protect critical sections, similar to a mutex

```
semaphore s;  
  
wait(s);  
    // Critical Section  
signal(s);
```

- each semaphore maintains a queue of processes blocked on the semaphore
- when a semaphore is locked by a thread, it can be unlocked by **any** thread
- as opposed to mutex, where a locking/unlocking must be done by the same thread

# Binary semaphore

- special type of semaphore with value either 0 or 1
- possible pseudocode implementation:

```
void wait(int s) {  
    while (s == 0) {;}  
    s = 0;  
}
```

```
void signal(s) {  
    s = 1;  
}
```

where the bodies are executed **atomically**

- **atomic operation:**
  - is an operation that appears to execute instantaneously with respect to the rest of the system
  - eg. cannot be interrupted by signals, threads, interrupts, ...



# Counting semaphore

- aka. **general semaphore**, or just **semaphore**
- represents an integer value  $S$
- $S > 0$ : value of  $S$  is the number of processes/threads that can issue a wait and immediately continue to execute
- $S = 0$ : eg. all resources are busy, the calling process/thread must wait
- $S < 0$ : some implementations might do this,  $\text{abs}(S)$  then represents the number of processes that are waiting to be unblocked
- possible pseudocode implementation:

```
void wait(int s) {  
    while (s <= 0) {}  
    s --;  
}
```

```
void signal(s) {  
    s ++;  
}
```

where the bodies are executed **atomically**

# Counting semaphore

- aka. **general semaphore**, or just **semaphore**
- represents an integer value  $S$

$S > 0$	value of $S$ is the number of threads that can issue <code>wait()</code> without blocking
$S == 0$	all resources are busy, any thread that calls <code>wait()</code> will block
$S < 0$	some implementations allow this... $ S $ then represents the number of threads that are blocked on <code>wait()</code>

- possible pseudocode implementation:

```
void wait(int s) {  
    while (s <= 0) {}  
    s --;  
}
```

```
void signal(s) {  
    s ++;  
}
```

where the bodies are executed **atomically**

# Semaphores vs. condition variables

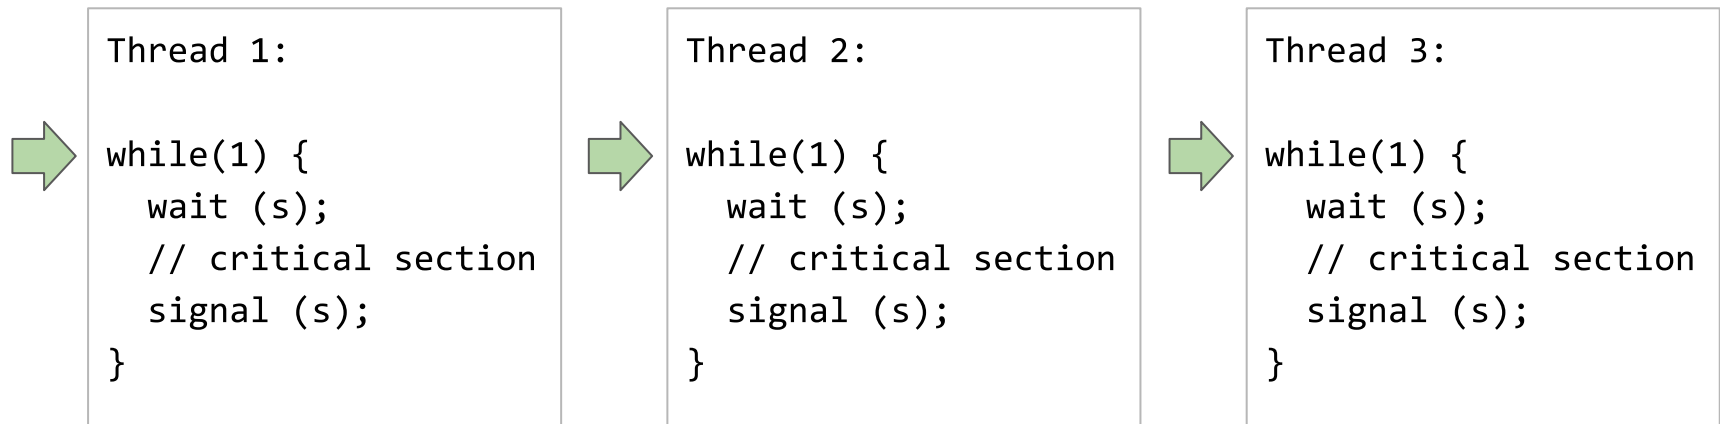
- `signal()` compared to `cv_signal()`:
  - `cv_signal()` is lost (has no effect) if no thread is waiting
  - `signal()` increments the semaphore always, and possibly wakes up a thread
- `wait()` compared to `cv_wait()`:
  - `cv_wait()` does not check the condition, it always blocks
  - `wait()` checks the value of the semaphore, and may or may not block

# POSIX semaphores

- `int sem_init (sem_t *sem, int pshared, unsigned int value)`  
initializes semaphore to `value`, `pshared=0/1` shared by threads/processes
- `int sem_destroy (sem_t * sem)`  
destroys the semaphore, fails if some threads are waiting on it
- `int sem_wait (sem_t * sem)`  
suspends the calling thread until the semaphore is non-zero, then atomically decreases the semaphore count
- `int sem_post (sem_t * sem)`  
atomically increases the semaphore, never blocks, may unblock blocked threads, safe to use in signal handlers in Linux on 486+ hardware
- `int sem_getvalue (sem_t * sem, int * sval)`  
returns the value of semaphore via `sval`
- `int sem_trywait (sem_t * sem)`  
non blocking version of `sem_wait()`

# Example with semaphore = 2

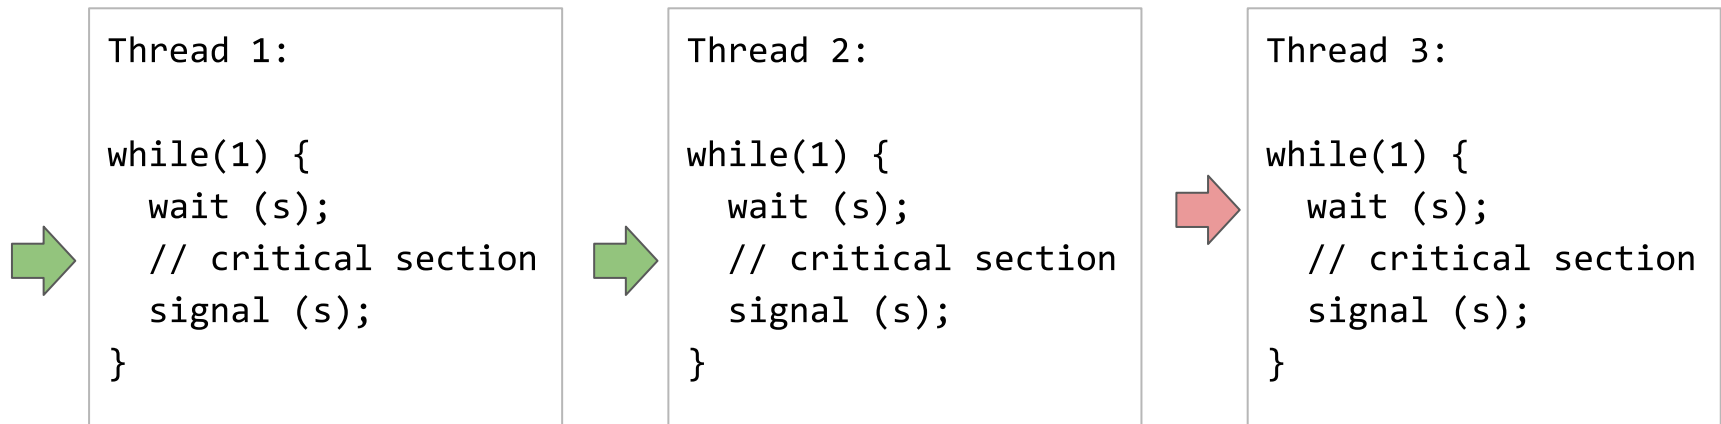
```
sem_t s;  
...  
sem_init( & s, 0, 2); // initialize semaphore to 2
```



- suppose 3 threads try to enter their CSs protected by a semaphore

# Example with semaphore = 2

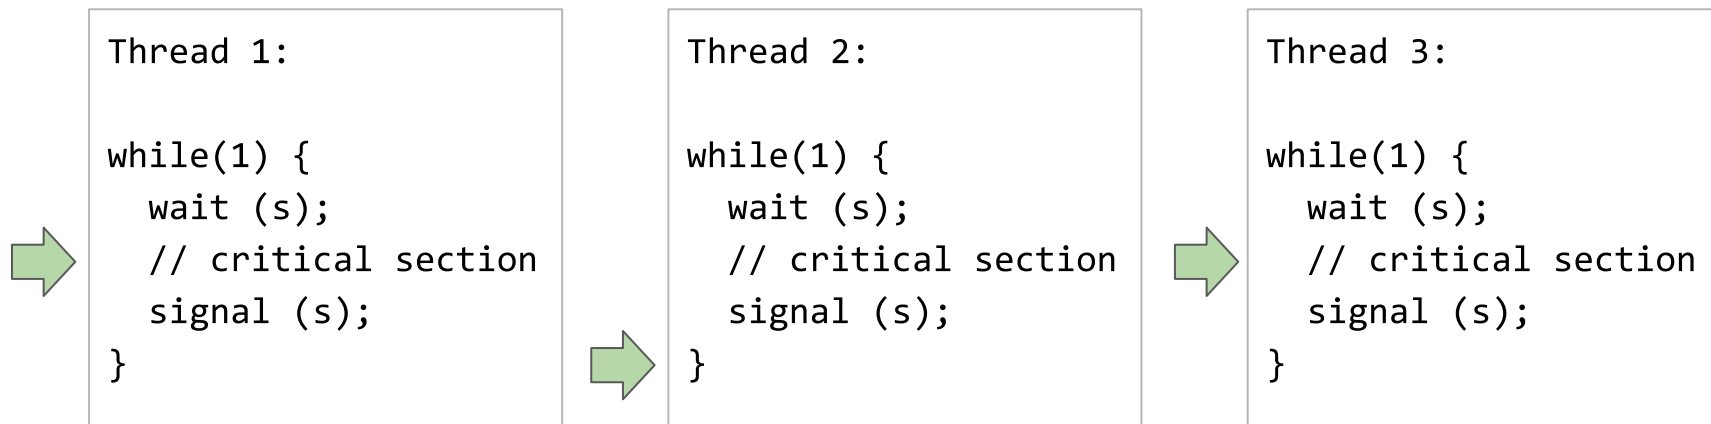
```
sem_t s;  
...  
sem_init( & s, 0, 2); // initialize semaphore to 2
```



- two threads will enter their CS simultaneously
- the other thread will be blocked

# Example with semaphore = 2

```
sem_t s;  
...  
sem_init( & s, 0, 2); // initialize semaphore to 2
```



- as soon as one thread leaves CS, the last thread will be allowed to enter its CS

# Improved semaphore implementation, with a queue & blocking

- S->list is a list of processes/threads
- a process can block() itself
- a process can be unblocked by wakeup()
- getpid() returns process or thread ID

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

```
void wait(semaphore *S) {  
    S->value --;  
    if (S->value < 0) {  
        S->list->push(getpid());  
        block();  
    }  
}
```

```
void signal(semaphore *S) {  
    S->value ++;  
    if (S->value <= 0) {  
        pid = S->list->pop();  
        wakeup(pid);  
    }  
}
```



# It is tricky

- one subtle error and everything comes to a grinding halt
- concurrent programming can be more difficult than coding in assembly
- any error with semaphores will potentially result in race conditions, deadlocks, and other forms of unpredictable and irreproducible behaviour
- for example:

Violate mutual exclusivity:

```
signal(sem);  
...  
critical section  
...  
wait(sem);
```

Deadlock:

```
lock(mutex);  
...  
critical section  
...  
lock(mutex);
```

# Another common problem with semaphores

- Scenario: managing a pool of  $N$  resources
  - semaphore  $S$  is used to keep track of the # of available resources
  - initialization: `sem_init( $S, 0, N$ )`
  - each process may need  $K_i \leq N$  resources at a time
  - you might be tempted to accomplish this by  $K_i$  consecutive invocations of `wait( $S$ )`
  
- Can you see the problem?
  - Example:  $N = 10$ , thread 1 needs 7 resources and thread 2 needs 6 resources
  - depending on scheduling, we may get a **deadlock**

# Another common problem with semaphores

## Thread 1:

```
for(i=0; i<7; i++)  
    wait(S);  
  
// critical section  
  
for(i=0; i<7; i++)  
    signal(S);  
...
```

## Thread 2:

```
for(i=0; i<6; i++)  
    wait(S);  
  
// critical section  
  
for(i=0; i<6; i++)  
    signal(S);  
...
```

- order of operations
- thread 1 requests 6 resources, then scheduler switches to thread 2
- thread 2 requests 4 resources, exhausting all available resources
- both threads are stuck → **deadlock**

# Extended semaphore - possible implementation

```
void wait(semaphore *S, int k) {  
    S->value -= k;  
    if (S->value < 0) {  
        S->list->push(getpid());  
        block();  
    }  
}
```

```
void signal(semaphore *S, int k) {  
    S->value += k;  
    if (S->value >= 0) {  
        pid = S->list->pop();  
        wakeup(pid);  
    }  
}
```

**Thread 1:**

```
wait(S, 7);
```

```
// critical section
```

```
signal(S, 7);
```

**Thread 2:**

```
wait(S, 6);
```

```
// critical section
```

```
signal(S, 6);
```

# Summary

---

- producer-consumer problem
- condition variables
- semaphores (binary, counting, extended)

Reference: 2.3 (Modern Operating Systems)

5.5, 5.8, 6.5 (Operating System Concepts)

Additional reading:

- The Little Book of Semaphores

<http://greenteapress.com/wp/semaphores/>

- Race conditions are not a problem among processes, only among threads. True or False?
- What is the main difference between `pthread_cond_signal()` for mutex and `sem_post()` for semaphore?
- A mutex is identical to binary semaphore. True or False?
- What does the value of the semaphore tell you?
- Define atomic operation.

# Dining philosophers with semaphores

```
#define N 5 /* number of philosophers */
#define THINKING 0 /* philosopher is thinking */
#define HUNGRY 1 /* philosopher is trying to get forks */
#define EATING 2 /* philosopher is eating */

int state[N]; /* each philosophers is either THINKING, HUNGRY or EATING, initialized to THINKING */
sem_t cs_m; /* semaphore for the critical section, shared by all philosophers, initialized to 1 */
sem_t p_m[N]; /* one semaphore per philosopher, initialized to 0 */

int left(i) { return (i+N-1) % N; }
int right(i) { return (i+1) % N; }

void philosopher(int i) /* to be executed by different threads */
{
    while (1) {
        think( ); /* philosopher is thinking */
        take_forks(i); /* acquire two forks or block */
        eat( ); /* yum-yum, spaghetti */
        put_forks(i); /* put both forks back on table */
    }
}
```

# Dining philosophers with semaphores

```
void take_forks(int i) {
    down(& cs_m);           /* enter critical region */
    state[i] = HUNGRY;      /* record fact that philosopher i is hungry */
    test(i);                /* try to acquire 2 forks */
    up(& cs_m);              /* exit critical region */
    down(& p_m[i]);          /* block if forks were not acquired */
}

void put_forks(i) {         /* i: philosopher number, from 0 to N-1 */
    down(& cs_m);           /* enter critical region */
    state[i] = THINKING;    /* philosopher has finished eating */
    test(left(i));          /* see if left neighbor can now eat */
    test(right(i));         /* see if right neighbor can now eat */
    up(& cs_m);              /* exit critical region */
}

void test(i) {              /* i: philosopher number, from 0 to N-1 */
    if (state[i] == HUNGRY && state[left(i)] != EATING && state[right(i)] != EATING) {
        state[i] = EATING; /* only start eating if hungry, and neighbors not eating */
        up(& p_m[i]);
    }
}
```



# Questions?