

CPSC 457

More synchronization mechanisms

Contains slides from Mea Wang, Andrew Tanenbaum and Herbert Bos

- monitors, spinlocks, event flags, mailboxes/messages
- mechanisms/algorithms for achieving race-free solutions
 - disabling interrupts
 - lock variables
 - strict alternation
 - Peterson's algorithm
 - synchronization hardware

Monitors

- a monitor is a higher-level construct compared to mutexes and semaphores
- a monitor is a **programming language construct** that controls access to shared data
 - synchronization code automatically added by a compiler, then enforced at runtime
 - implemented in Concurrent Pascal, C#, D, Modula-3, Java, Ruby, Python, ...
 - can be somewhat emulated in C++ (classes)
and even to some extent in C (struct + functions or function pointers)

Monitors

- a monitor is a **module** that encapsulates
 - shared data structures
 - and methods that operate on these shared data structures
 - synchronization is done between concurrent method invocations
- data in monitors can only be accessed via the published methods
- a properly implemented monitor is virtually impossible to use in a wrong way

- another way to look at a monitor:
 - a thread-safe class/object
 - automatic mutual exclusion on every method (via built-in mutex)
 - can include condition variables for signalling conditions

Monitors

- implement automatic mutual exclusion
- only one thread can execute any monitor method at any time (per instance)
- all other threads would block if they tried
- monitors have their own queues
- Example:

```
Monitor counter {  
    int c = 0;  
    void incr() {  
        c = c + 1;  
    }  
    void decr() {  
        c = c - 1;  
    }  
}
```

- ← calling `incr()` or `decr()` from multiple threads would allow one thread in, the rest would block
- ← think of all bodies of all methods as being critical sections, protected by one mutex
- ← in C++ you can emulate this by making a private mutex, and locking it at the beginning of every method

Monitors with condition variables

- monitors can have their own condition variables
- CVs declared as part of the module
- only accessible from within the module (private)
- similar to `pthread_cond_t`

```
Monitor counter {  
    ...  
    condition c1, c2;  
    proc() {  
        c1.wait();  
        c2.signal();  
    }  
}
```

Monitors vs. semaphores and mutexes

- once a monitor is correctly programmed, access to the protected resource is correct for accessing from all threads
- with semaphores or mutexes, resource access is correct only if all threads that access the resource are programmed correctly
- programming with monitors → you test/debug the monitor
- programming with mutexes/semaphores → you test/debug all code using them

Spinlocks

- another synchronization mechanism
- lightweight alternative to mutex
- implemented using busy waiting loops
- usually implemented in assembly, using **atomic operations**
- very efficient if you know the wait time will be very short
 - because no re-scheduling required
 - only make sense on multi-core/CPU systems
- often used inside kernels

atomic operation - an operation that appears to execute instantaneously w.r.t. to the rest of the system, i.e. it cannot be *interrupted* by anything else

Spinlocks

```
mutex m;
```

```
mutex_lock( & m);
```

```
    /* critical section */
```

```
mutex_unlock( & m);
```

```
spinlock s;
```

```
spin_lock( & s);
```

```
    /* (very short) critical section */
```

```
spin_unlock( & s);
```

Spinlock in x86 - <https://en.wikipedia.org/wiki/Spinlock>

```
locked:                ; The lock variable. 1 = locked, 0 = unlocked.
    dd                0 ; initialized to 'unlocked'

spin_lock:             ; procedure to lock the mutex
    mov     eax, 1      ; set the EAX register to 1
    xchg    eax, [locked] ; atomically swap EAX with the lock variable
                                ; eax = old locked, new locked = 1
    test    eax, eax    ; test whether old locked == 0
    jnz     spin_lock   ; keep spinning until old locked == 0
    ret                                ; the lock was acquired, we are done (locked = 1!!!)

spin_unlock:           ; procedure to unlock the mutex
    mov     eax, 0      ; set EAX register to 0.
    xchg    eax, [locked] ; atomically set locked = 0
    ret                                ; lock has been released
```

Spinlocks in pthreads

```
#include <pthread.h>
```

```
int  pthread_spin_init(pthread_spinlock_t * lock, int pshared);
```

```
int  pthread_spin_destroy(pthread_spinlock_t * lock);
```

```
int  pthread_spin_lock(pthread_spinlock_t * lock);
```

```
int  pthread_spin_trylock(pthread_spinlock_t * lock);
```

```
int  pthread_spin_unlock(pthread_spinlock_t * lock);
```

Notice the similarity to mutex APIs - making it trivial to switch between mutexes and spinlocks.

More synchronization mechanisms

■ **event flags:**

- a memory word with N bits
 - different events may be associated with different bits in a flag
- operations:
 - set flag
 - clear flag
 - wait for 1 flag
 - wait for any flag
 - wait for all flags

More synchronization mechanisms

■ message passing

- processes send each other **messages**
- messages can contain arbitrary data
- delivered messages can be queued up in **mailboxes**
- processes can check contents of mailboxes, take messages out, or wait for messages
- common implementation is MPI (message passing interface)
 - very popular in HPC (high-performance-computing)
 - many good tutorials available (google "MPI tutorial")

Priority inversion

- Scenario: Assume we have three processes, L, M, and H, whose priorities follow the order $L < M < H$. Assume that process H requires resource R, which is currently being accessed by process L. While H is waiting for L to finish using resource R, M becomes runnable, thereby preempting process L. Now, H has to wait for both M and L to finish. Effectively, H has the lowest priority in this execution.
- This problem is known as [priority inversion](#).
- Possible solution: [priority-inheritance](#)
 - As soon as H requests resource R, the process P holding resource R automatically inherits the priority of H if the P has lower priority. Once P releases the resource, its original priority is restored.
 - As a result, we would have the following execution order in the above scenario: L, H, M.

Race conditions with separate programs

- Scenario - naive print server
 - if a program needs to print something, the program 'submits' the print job by writing a PDF files to the directory /printJobs
 - a print server monitors /printJobs for jobs to print
 - eg. by scanning the directory for *.pdf every N seconds
 - when the server finds a PDF file, it prints it and removes the file
 - if it does not find any PDF files, it sleeps for N seconds
- Can you spot a problem with this?

Race conditions with separate programs

- Scenario - naive print server
 - if a program needs to print something, the program 'submits' the print job by writing a PDF files to the directory /printJobs
 - a print server monitors /printJobs for jobs to print
 - eg. by scanning the directory for *.pdf every N seconds
 - when the server finds a PDF file, it prints it and removes the file
 - if it does not find any PDF files, it sleeps for n seconds
- Can you spot a problem with this?
 - print server might print an incomplete PDF file
- How do we fix this?
 - we must prevent incomplete PDF files appearing in the directory
 - we need to make sure that when the PDF file appears in the directory, it is complete

Race conditions with separate programs

- can you design a solution?
- what assumptions can we make?
- some filesystems support file/directory locking mechanisms, but not all, so we should not rely on those
- however, nearly all filesystems support at least 2 atomic operations:
 - file creation
 - file rename

Race conditions with separate programs

- solution - we modify the programs that print, while print server needs no modification
- when a program needs to print:
 1. it **creates** a temporary file in `/printJobs` directory, with a random filename, eg. `job-xxxxxx.tmp`, where 'xxxxxx' is a unique number
 2. it writes the output to this temporary file
 3. when done, it **renames** `job-xxxxxx.tmp` to `job-yyyyyy.pdf` where 'yyyyyy' is a another unique number
- creating random file without a race condition:
 - there is a UNIX utility function `mkstemp()`, although you could make your own easily
- renaming to a random filename – no utility, need to make our own, eg.
loop:
 generate random string yyyyyy
 call `rename(job-xxxxxx.tmp, job-yyyyyy.pdf)` // atomic system call
 if successful, break out of loop

creating a file with

```
open (pathname,  
      O_CREAT | O_EXCL,  
      0644)
```

is atomic

Readers/writers problem

- scenario:
 - a single resource is shared among several threads
 - some threads only read the resource, others only write it
 - the resource supports **multiple concurrent readers**, but only a **single writer**
 - also called **shared-exclusive lock**, or **reader-writer lock**
- can we use semaphores to implement this efficiently?
 - note — pthreads has `pthread_rwlock_*`() functions for just this purpose
- but we can design our own solution using three variables:
 - `int readCount` - contains number of readers reading
 - `semaphore cs` - used as a mutex to control access to critical sections
 - `semaphore w_only` - used as flag representing whether write is available

Readers/writers problem with semaphores

```
int readcount = 0;
sem_t cs = 1;
sem_t w_only = 1;

writer {
    wait(w_only); // lock out readers
    do_writing();
    signal(w_only); // allow readers
}
```

```
reader {
    wait(cs); // enter critical section
    readCount ++; // add reader
    if (readCount == 1)
        wait(w_only); // wait for writers
    signal(cs); // exit critical section
    do_reading();
    wait(cs); // enter critical section
    readCount --; // remove reader
    if (readCount == 0)
        signal(w_only); // let writers write
    signal(cs); // exit critical section
}
```

Note: first reader locks w_only,
last reader releases w_only

Semaphore implemented with mutexes and cond. vars.

```
struct sem {  
    pthread_mutex_t mutex;  
    int count = 0;  
    pthread_cond_t cond;  
};
```

```
void down(sem * s) {  
    lock(s->mutex);  
    while(s->count == 0)  
        wait(s->mutex, s->cond);  
    s->count --;  
    unlock(s->mutex);  
}
```

```
void up(sem * s) {  
    lock(s->mutex);  
    s->count ++;  
    signal(s->cond);  
    unlock(s->mutex);  
}
```

Requirements for good race-free solution


Recall:

1. **Mutual exclusion:** No two processes/threads may be simultaneously inside their critical sections (CS).
2. **Progress:** No process/threads running outside its CS may block other processes/threads.
3. **Bounded waiting:** No process/thread should have to wait forever to enter its CS.
4. **Speed:** No assumptions may be made about the speed or the number of CPUs.

General structure of a process

```
while (1) {  
    entry section  
    critical section  
    exit section  
    non-critical section  
}
```

Achieving mutual exclusion

- blocking techniques:
 - mutex, semaphore, monitor, condition variables 
- non-blocking techniques:
 - disabling interrupts
 - lock variables
 - strict alternation
 - Peterson's algorithm
 - synchronization hardware
 - spinlocks

Disabling interrupts

The idea:

- Each process disables all interrupts just before entering its CS and re-enable them just before leaving the CS.
- Once a process has disabled interrupts, it can examine and update the shared memory without interventions from other processes.

Problems:

- What if a process never re-enables the interrupts?
- On a multi-CPU systems, disabling interrupts affects only one CPU.
- Sometimes used inside kernels, but even that is becoming problematic.

```
while(1) {  
    disable interrupts  
    critical section  
    enable interrupts  
    non-critical section  
}
```


Lock variables

- A single, shared (lock) variable, initialized to 0.
 - `lock == 0` : no process is in CS
 - `lock == 1` : a process is in CS
- A process can only enter its CS if `lock == 0`.
Otherwise, it must wait.
- Any problems?

```
while(1) {  
    while (lock == 1) {}  
    lock = 1;  
    critical section  
    lock = 0;  
    non-critical section  
}
```

Lock variables

- A single, shared (lock) variable, initialized to 0.
 - `lock == 0` : no process is in CS
 - `lock == 1` : a process is in CS
- A process can only enter its CS if `lock == 0`.
Otherwise, it must wait.
- **Problem: no mutual exclusion**
 - 1st thread gets past while...
 - context switch to 2nd thread
 - 2nd thread enters CS
 - context switch to 1st thread
 - 1st thread enters CS
 - both threads are in their CS

```
while(1) {  
    while (lock == 1) {}  
    lock = 1;  
    critical section  
    lock = 0;  
    non-critical section  
}
```

Lock variables

- A single, shared (lock) variable, initialized to 0.
 - 0 - no process is in CS
 - 1 - a process is in CS
- A process can only enter its CS if lock is 0.
Otherwise, it must wait.
- **Problem: no mutual exclusion**
- **Solution:**
 - make entering CS atomic operation
eg. using compare-and-swap
 - or ...

```
while(1) {  
    while (lock == 1) {}  
    lock = 1;  
    critical section  
    lock = 0;  
    non-critical section  
}
```

Strict alternation

- two processes alternate entering their critical sections
- global variable: `turn = 0`

Thread 1:

```
while(1) {  
    while (turn != 0) {}  
    critical section  
    turn = 1;  
    non-critical section  
}
```

Thread 2:

```
while(1) {  
    while (turn != 1) {}  
    critical section  
    turn = 0;  
    non-critical section  
}
```

- mutual exclusion is OK
- other problems?

Strict alternation

- two processes alternate entering their critical sections
- global variable: `turn = 0`

Thread 1:

```
while(1) {  
    while (turn != 0) {}  
    critical section  
    turn = 1;  
    non-critical section  
}
```

Thread 2:

```
while(1) {  
    while (turn != 1) {}  
    critical section  
    turn = 0;  
    non-critical section  
}
```

← spinlock

- **problem 1**: busy waiting (is it always a problem?)
- **problem 2**: only works for 2 processes
- **problem 3**: no progress - faster process is blocked by slower process not in its CS

Peterson's algorithm

- software solution for 2 processes (can be extended to N processes)
 - assumes atomicity, visibility & ordering
 - eg. may fail on some CPUs with out-of-order execution, or memory reordering
- shared integer variable **turn** - indicating whose turn it is
- shared array **flag[2]** - indicating who is interested in entering CS, initialized to 0

Process 0:

```
while(1) {  
    flag[0] = TRUE;  
    turn = 1;  
    while (flag[1] && turn == 1) {}  
        critical_section  
    flag[0] = FALSE;  
  
    non-critical-section  
}
```

Process 1:

```
while(1) {  
    flag[1] = TRUE;  
    turn = 0;  
    while (flag[0] && turn == 0) {}  
        critical_section  
    flag[1] = FALSE;  
  
    non-critical-section  
}
```

Peterson's algorithm

- software solution for 2 processes (can be extended to N processes)
 - assumes atomicity, visibility & ordering
 - eg. may fail on some CPUs with out-of-order execution, or memory reordering
- shared integer variable **turn** - indicating whose turn it is
- shared array **flag[2]** - indicating who is interested in entering CS, initialized to 0

Process 0:

```
while(1) {
    flag[0] = TRUE;
    turn = 1;
    while (flag[1] && turn == 1)
        critical_section
    flag[0] = FALSE;

    non-critical-section
}
```

Process 1:

```
{
    flag[1] = TRUE;
    turn = 0;
    while (flag[0] && turn == 0) {}
    critical_section
    flag[1] = FALSE;

    non-critical-section
}
```

mutual exclusion ✓

progress ✓

bounded waiting ✓

Synchronization hardware

- race conditions are prevented by ensuring that critical sections are protected by locks:
 - a process must acquire a lock before entering a CS
 - a process releases the lock when it exits the CS
- many modern computer systems provide special hardware instructions that implement useful atomic operations
 - these can be used to create atomic locking and unlocking mechanisms
 - e.g. compare-and-swap, test-and-set, swap

Compare-and-swap (CAS)

- atomic operation used for synchronization
- supported by most CPUs today, eg. CMPXCHG on Intel
- general algorithm:

- compare contents of memory to val1
- if they are the same, change the memory to val2
- return the old contents of memory

} must be atomic !

- pseudo-code:

```
int cas(int * mem, int val1, int val2)
{
    int old = * mem;
    if (old == val1) * mem = val2;
    return old;
}
```

} must be atomic !

```
int p = 0;

while(1) {
    while( cas(&p,0,1) == 1){}
    // critical section
    p = 0;
    // non-critical section
}
```

Compare-and-swap in GCC 4.4+

- gcc provides a number of atomic operations, including CAS
- `type __sync_val_compare_and_swap (type *ptr, type oldval, type newval)`
 - atomic compare and swap
 - if the current value of *ptr is oldval, then write newval into *ptr
 - returns the original value of *ptr
- `bool __sync_bool_compare_and_swap (type *ptr, type oldval, type newval)`
 - same as above, but return true if newval was written
- type - can be any 8, 16, 32 and 64 bit integers or pointers

Spinlock using compare-and-swap

```
void spin_lock(volatile int *p)
{
    while(!__sync_bool_compare_and_swap(p, 0, 1)) { /* ??? */ }
}
```

```
void spin_unlock(volatile int *p)
{
    *p = 0;
}
```

Spinlock using compare-and-swap

```
void spin_lock(volatile int *p) {
    while(!__sync_bool_compare_and_swap(p, 0, 1)) {

        ;                // option 1 - CPU/bus very busy
        sched_yield();    // option 2 - reschedule thread, lot of CPU overhead
        while(*p) {};     // option 3 - lot of CPU overhead, less bus overhead
        while(*p) _mm_pause(); // option 4 - less bus overhead with multiple CPUs
                                // mm_pause is an intrinsic, delaying
                                // execution of the next instruction
                                // by a small amount
    }
}

void spin_unlock(volatile int *p) {
    *p = 0;
}
```

Test-and-set

- a specialized version of compare-and-swap
 - old hardware had test-and-set
 - newer hardware uses more generalized compare-and-swap

- general algorithm

- remember contents of memory
 - set memory to true
 - return the old contents of memory
- } atomic !

- pseudo-code:

```
int tas(int* mem)
{
    int old = *mem;
    *mem = TRUE;
    return old;
}
```

} atomic !

```
int p = 0;

while(1) {
    while( tas(&p) ) {}
    critical section
    p = 0;
    non-critical section
}
```

Swap

- another atomic operation that can be used for synchronization
- general algorithm
 - atomically swap contents of two memory locations
- pseudo-code:

```
void swap(int* a, int* b)
{
    int tmp = * a;
    * a = * b;
    * b = tmp;
}
```

} atomic !

```
int lock = 0;

while(1) {
    int key = TRUE;
    while( key == TRUE ) {
        swap( &lock, &key);
    }
    critical section
    lock = FALSE;
    non-critical section
}
```

Bounded waiting with synchronization hardware

- when used correctly, the atomic operations, such as compare-and-swap, test-and-set, swap can be used to achieve mutual exclusion, progress and speed
- but they are too low level to achieve bounded waiting, especially for more than 2 processes
- bounded waiting can be 'added', eg. via two shared variables

```
// a CS lock
int lock;
// which processes want to enter CS
int waiting[n];
```

```
int id; // process's ID, starting at 0

while (1) {
    waiting[id] = true;
    while (waiting[id]
           && testAndSet(&lock)) {}
    waiting[id] = false;

    /* critical section */

    j = (id + 1) % n;
    while ((j != id) && !waiting[j])
        j = (j + 1) % n;
    if (j == id)
        lock = false;
    else
        waiting[j] = false;

    /* remainder section */
}
```

Bounded waiting with synchronization hardware

indicate interest to enter CS

if CS is locked, we wait until someone else gives us a turn

clear indicator

find next process trying to enter CS

if nobody else waiting, release lock

if we found someone else waiting, let them enter CS, but without releasing lock

```
int id; // process's ID, starting at 0

while (1) {
    waiting[id] = true;
    while (waiting[id]
           && testAndSet(&lock)) {};
    waiting[id] = false;

    /* critical section */

    j = (id + 1) % n;
    while ((j != id) && !waiting[j])
        j = (j + 1) % n;
    if (j == id)
        lock = false;
    else
        waiting[j] = false;

    /* remainder section */
}
```


Synchronization hardware

- can be used to implement mutual exclusion, progress, speed and even bounded waiting
- drawbacks:
 - busy-waiting (spinlocks)
 - extra coding (eg. to add bounded waiting)
- advantage:
 - avoids system calls
 - can be more efficient if expected wait time is short
 - only makes sense on multi-CPU/core systems

Summary

- monitors, spinlocks, event flags, mailboxes/messages
- lock-free solutions:
 - disabling interrupts, lock variables, strict alternation, Peterson's algorithm
 - synchronization hardware: compare-and-swap, test-and-set, swap
- free book !!!
The Little Book of Semaphores - <http://greenteapress.com/wp/semaphores/>

Reference: 2.3 (Modern Operating Systems)

5.5, 5.8, 6.5 (Operating System Concepts)

- Which one of the following achieves mutual exclusion but violates the “progress” requirement?
 - a) Disabling Interrupts
 - b) Lock Variables
 - c) Strict Alternation
 - d) Peterson’s Solution

Questions?